# Report by Danil Ginzburg

BS17-DS-01

#### Transforming images to the same size

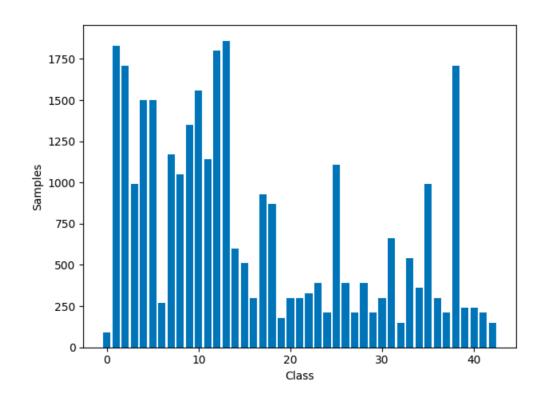
The transformation is done by padding image to obtain square shape and then resizing it to particular state (typically it is 30 \* 30). Padding is done via replication of neighbour pixels.

Here is an example of rectangular image (182, 67, 3) and its transformation to (30, 30, 3):



#### Image frequencies

After splitting data into training and validation sets, 80% of images went to training set and the rest went to validation set. Here is bar representing the number of samples for each class:



### Augmentation

In order to equalise the number of samples per class we are using Augmentation. Here is two types of augmentation I applied for images: **Add light** and **Contrast**. Since all signs are already centralised and since a flipped sign is not a sign anymore, thus such techniques as rotate, flip, etc. are meaningless. Therefore, I find Add light and Contrast the most useful augmentation techniques.

Here are the examples of them:

I. Add light. Images — Original; gamma=2; gamma=3; gamma=4.



II. **Contrast image**. Images — Original; contrast=10; contrast=20; contrast=30.



#### Evaluation

On the right side is the result of calling		precision	recall	f1-score	support
classification_report function from sklearn.metrics.	0	0.52	0.58	0.55	60
	1	0.67	0.64	0.66	450
	10	0.86	0.93	0.89	390
<pre>clf = RandomForestClassifier(n_estimators=30, max_depth=50)</pre>	11	0.86	0.82	0.84	360
<pre>clf.fit(X, y)</pre>	12	0.98	0.90	0.94	630
<pre>y_pred = clf.predict(X_test)</pre>	13	0.92	0.83	0.87	360
<pre>print(classification_report(y_test, y_pred))</pre>	14	0.90	0.98	0.94	150
print(ctassification_report(y_test, y_prea/)	15	0.66	0.78	0.72	60
	16	0.57	0.27	0.36	30
On the left side we see the number of class	17	0.97	0.98	0.98	120
	18	0.77	0.78	0.78	270
(0-42) and its precision, recall, f1-score, etc.	19	0.87	0.66	0.75	90
	2	0.55	0.72	0.62	360
The overall accuracy is appeared to be 75% with	20 21	0.43	0.65	0.52	60
		0.00	0.00	0.00	0
image shape = $(30, 30)$ and on augmented data.	22	0.79	0.92	0.85	90
Although it can be increased (see «Experiments»	23	0.81	0.45	0.58	210
subtitle below).	24	0.48	0.50	0.49	30
Subtitio bolow).	25	0.91	0.81	0.86	360
	26	0.46	0.52	0.48	60
	27	0.60	0.90	0.72	30
	28	0.87	0.78	0.82	150
	29	0.77	0.60	0.68	120
	3	0.65	0.63	0.64	390
	30	0.63	0.52	0.57	120
	31	0.50	0.70	0.58	90
	32	0.52	0.20	0.29	60
	33	0.97	0.99	0.98	149
	34	0.83	0.97	0.89	91
	35	0.98	0.91	0.95	239

36

37

38

0.86

1.00

0.97

0.98

0.92

0.95

0.92

0.96

91

89

481

39	0.00	0.00	0.00	0
4	0.72	0.75	0.73	390
40	0.73	1.00	0.85	30
41	1.00	0.60	0.75	60
42	0.00	0.00	0.00	0
5	0.39	0.39	0.39	360
6	0.54	0.93	0.68	90
7	0.40	0.47	0.43	210
8	0.66	0.62	0.64	390
9	0.97	0.70	0.81	240
accuracy			0.75	8010
macro avg	0.69	0.68	0.67	8010
weighted avg	0.78	0.75	0.76	8010

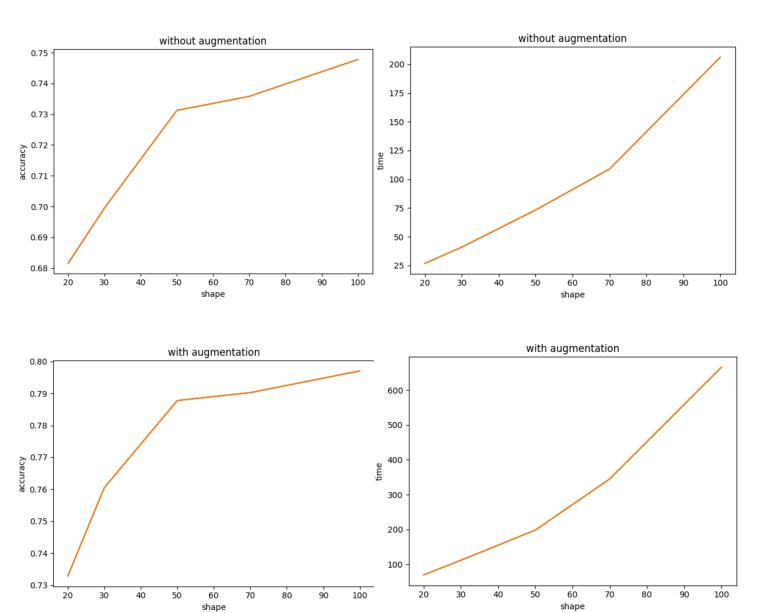
Some images has been misclassified and here are some examples of those:



The obvious reason they have been misclassified is because they are blurred or too dark.

## **Experiments**

Here is result plots of doing training on **non augmented** data with different shapes - (20, 20); (30, 30); (50, 50); (70, 70); (100, 100):



And here is result plots of doing the same training, but on **augmented** data: As we observe, the larger the size, the better the accuracy. So choosing appropriate size is only about tradeoff between time to train and accuracy.

Augmentation takes time itself, and training on augmentation data is much longer (in this case 3 times longer in general). However, it definitely worth it, since we managed to increase accuracy by 5% in general.